

# Data\_Analysis

August 28, 2017

## 1 Amazon Fine Food Reviews: EDA

### 1.1 Introduction

Data Fields Explanation

The Amazon Fine Food Reviews dataset consists of 568,454 food reviews. This dataset consists of a single CSV file, Reviews.csv. The columns in the table are:

```
In [63]: #from IPython.display import Image
         #Image(filename='AmazonReview.png')
```

### 1.2 Objective

Analysing the data & plot the required graphs to show that these conclusions are true:

Note: This notebook is highly inspired from the Exploratory visualization of Amazon fine food reviews by Rob Castellano.

### 1.3 Loading the Data

```
In [64]: #Let's import pandas to read the csv file.
import pandas as pd
df = pd.read_csv("Reviews.csv")
```

```
In [65]: #Printing first 5 columns from our data frame
df.head()
```

```
Out [65]:
```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham	"M. Wassir"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	5	1303862400	
1	0	0	1	1346976000	
2	1	1	4	1219017600	
3	3	3	2	1307923200	

4 0 0 5 1350777600

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...
3	Cough Medicine	If you are looking for the secret ingredient i...
4	Great taffy	Great taffy at a great price. There was a wid...

```
In [66]: #Observing the lables of each column
print(df.keys())
```

```
Index([u'Id', u'ProductId', u'UserId', u'ProfileName', u'HelpfulnessNumerator',
      u'HelpfulnessDenominator', u'Score', u'Time', u'Summary', u'Text'],
      dtype='object')
```

```
In [67]: #Observing the shape of our data frame.
df.shape
# Note: We have 10 features and 568454 data points.
```

```
Out [67]: (568454, 10)
```

```
In [68]: #Lets check for missing values
df.info()
#Observe that there are some missing values in 'PROFILENAME' & 'SUMMARY' column.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 568454 entries, 0 to 568453
Data columns (total 10 columns):
Id                568454 non-null int64
ProductId         568454 non-null object
UserId           568454 non-null object
ProfileName       568438 non-null object
HelpfulnessNumerator  568454 non-null int64
HelpfulnessDenominator  568454 non-null int64
Score            568454 non-null int64
Time             568454 non-null int64
Summary          568428 non-null object
Text             568454 non-null object
dtypes: int64(5), object(5)
memory usage: 43.4+ MB
```

```
In [69]: df.describe()
#Observe that more than 75% of our data is belonging to positive /
#class, i.e. we have imbalanced dataset.
```

```
Out [69]:
```

	Id	HelpfulnessNumerator	HelpfulnessDenominator	\
count	568454.000000	568454.000000	568454.000000	

mean	284227.500000	1.743817	2.22881
std	164098.679298	7.636513	8.28974
min	1.000000	0.000000	0.00000
25%	142114.250000	0.000000	0.00000
50%	284227.500000	0.000000	1.00000
75%	426340.750000	2.000000	2.00000
max	568454.000000	866.000000	923.00000

	Score	Time
count	568454.000000	5.684540e+05
mean	4.183199	1.296257e+09
std	1.310436	4.804331e+07
min	1.000000	9.393408e+08
25%	4.000000	1.271290e+09
50%	5.000000	1.311120e+09
75%	5.000000	1.332720e+09
max	5.000000	1.351210e+09


```
In [70]: #Lets do the value count on 'Scores'.
df.Score.value_counts()
```

```
Out[70]: 5    363122
         4    80655
         1    52268
         3    42640
         2    29769
         Name: Score, dtype: int64
```

## 1.4 Exploratory Data Analysis

Till now we saw that 5-star reviews constitute a large proportion (64%) of all reviews. The next most prevalent rating is 4-stars(14%), followed by 1-star (9%), 3-star (8%), and finally 2-star reviews (5%). Note that we have 10 features and 568454 data points. There are some missing values in 'PROFILENAME' & 'SUMMARY' column. More than 75% of our data is belonging to positive class(Score=4,5), i.e. we have imbalanced dataset.

```
In [71]: #Importing Seaborn and Matplotlib for graphical effects.
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure()
sns.countplot(x='Score', data=df, palette='RdBu')
plt.xlabel('Score (Rating)')
plt.show()
```



## 1.5 Creating a new dataframe

```
In [72]: #copying the original dataframe to 'temp_df'.
temp_df = df[['UserId','HelpfulnessNumerator','HelpfulnessDenominator', 'Summary', 'T

#Adding new features to dataframe.
temp_df["Sentiment"] = temp_df["Score"].apply(lambda score: "positive" if score > 3 e
                                                ("negative" if score < 3 else "not defin
temp_df["Usefulness"] = (temp_df["HelpfulnessNumerator"]/temp_df["HelpfulnessDenominator"]
(lambda n: ">75%" if n > 0.75 else "<25%" if n < 0.25 else "25-75%" if n >= 0.25 and
                                                n <= 0.75 else

temp_df.loc[temp_df.HelpfulnessDenominator == 0, 'Usefulness'] = ["useless"]
# Removing all rows where 'Score' is equal to 3
#temp_df = temp_df[temp_df.Score != 3]
#Lets now observe the shape of our new dataframe.
temp_df.shape
```

Out[72]: (568454, 8)

```
In [73]: temp_df.describe()
```

```
Out[73]:
```

	HelpfulnessNumerator	HelpfulnessDenominator	Score
count	568454.000000	568454.000000	568454.000000

mean	1.743817	2.22881	4.183199
std	7.636513	8.28974	1.310436
min	0.000000	0.00000	1.000000
25%	0.000000	0.00000	4.000000
50%	0.000000	1.00000	5.000000
75%	2.000000	2.00000	5.000000
max	866.000000	923.00000	5.000000

```
In [74]: temp_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 568454 entries, 0 to 568453
Data columns (total 8 columns):
UserId                568454 non-null object
HelpfulnessNumerator  568454 non-null int64
HelpfulnessDenominator 568454 non-null int64
Summary               568428 non-null object
Text                  568454 non-null object
Score                 568454 non-null int64
Sentiment             568454 non-null object
Usefulness            568454 non-null object
dtypes: int64(3), object(5)
memory usage: 34.7+ MB
```

```
In [75]: #Lets view the dataframe when Score=5
temp_df[temp_df.Score == 5].head(10)
```

```
Out [75]:
```

	UserId	HelpfulnessNumerator	HelpfulnessDenominator	\
0	A3SGXH7AUHU8GW	1	1	
4	A1UQRSCLF8GW1T	0	0	
6	A1SP2KVKFXXRU1	0	0	
7	A3JRGQVEQN31IQ	0	0	
8	A1MZY09TZKOBBI	1	1	
9	A21BT40VZCCYT4	0	0	
10	A3HDK07OWOQNK4	1	1	
11	A2725IB4YY9JEB	4	4	
14	A2MUGFV2TDQ47K	4	5	
15	A1CZX3CP8IKQIJ	4	5	

	Summary	\
0	Good Quality Dog Food	
4	Great taffy	
6	Great! Just as good as the expensive brands!	
7	Wonderful, tasty taffy	
8	Yay Barley	
9	Healthy Dog Food	
10	The Best Hot Sauce in the World	
11	My cats LOVE this "diet" food better than thei...	

```

14                      Strawberry Twizzlers - Yummy
15      Lots of twizzlers, just what you expect.

```

	Text	Score	Sentiment	\
0	I have bought several of the Vitality canned d...	5	positive	
4	Great taffy at a great price. There was a wid...	5	positive	
6	This saltwater taffy had great flavors and was...	5	positive	
7	This taffy is so good. It is very soft and ch...	5	positive	
8	Right now I'm mostly just sprouting this so my...	5	positive	
9	This is a very healthy dog food. Good for thei...	5	positive	
10	I don't know if it's the cactus or the tequila...	5	positive	
11	One of my boys needed to lose some weight and ...	5	positive	
14	The Strawberry Twizzlers are my guilty pleasur...	5	positive	
15	My daughter loves twizzlers and this shipment ...	5	positive	


	Usefulness
0	>75%
4	useless
6	useless
7	useless
8	>75%
9	useless
10	>75%
11	>75%
14	>75%
15	>75%

## 1.6 Positive reviews are very common

```

In [76]: sns.countplot(x='Sentiment', order=["positive", "negative"], data=temp_df, palette='R
plt.xlabel('Sentiment')
plt.show()

```



```
In [77]: temp_df.Sentiment.value_counts()
```

```
Out[77]: positive      443777  
         negative      82037  
         not defined   42640  
         Name: Sentiment, dtype: int64
```

Therefore we could conclude that the positive reviews are way more than the negative reviews.

## 1.7 Polpular words in Review

A look at the post popular words in positive (4-5 stars) and negative (1-2 stars) reviews shows that both positive and negative reviews share many popular words, such as "coffee", "taste", "flavor", "price", "good", and "product." The words "good", "great", "love", "favorite", and "find" are indicative of positive reviews, while negative reviews contain words such as "didn't" and "disappointed", but these distinguishing words appear less frequently than distinguishing words in positive reviews.

```
In [78]: pos = temp_df.loc[temp_df['Sentiment'] == 'positive']  
         pos = pos[0:25000]  
  
         neg = temp_df.loc[temp_df['Sentiment'] == 'negative']  
         neg = neg[0:25000]
```

```

In [79]: import nltk
        from nltk.corpus import stopwords
        from wordcloud import WordCloud
        import string
        import matplotlib.pyplot as plt

        def create_Word_Corpus(temp):
            words_corpus = ''
            for val in temp["Summary"]:
                text = str(val).lower()
                #text = text.translate(trantab)
                tokens = nltk.word_tokenize(text)
                tokens = [word for word in tokens if word not in stopwords.words('english')]
                for words in tokens:
                    words_corpus = words_corpus + words + ' '
            return words_corpus

        # Generate a word cloud image
        pos_wordcloud = WordCloud(width=900, height=500).generate(create_Word_Corpus(pos))
        neg_wordcloud = WordCloud(width=900, height=500).generate(create_Word_Corpus(neg))

```

C:\Users\Kanav\Anaconda2\lib\site-packages\ipykernel\_launcher.py:14: UnicodeWarning: Unicode e

```

In [80]: # Plot cloud
        def plot_Cloud(wordCloud):
            plt.figure( figsize=(20,10), facecolor='w')
            plt.imshow(wordCloud)
            plt.axis("off")
            plt.tight_layout(pad=0)
            plt.show()
            plt.savefig('wordclouds.png', facecolor='w', bbox_inches='tight')

In [81]: #Visuallizing popular positive words
        plot_Cloud(pos_wordcloud)

```





## 1.8 Helpfulness


### 1.8.1 How many reviews are helpful?

Among all reviews, almost half (50%) are not voted on at all. Among reviews that are voted on, helpful reviews(>75%) are the most common

```
In [83]: #Checking the value count for 'Usefulness'
temp_df.Usefulness.value_counts()
```

```
Out[83]: useless      270052
>75%      202836
25-75%     57286
<25%       38280
Name: Usefulness, dtype: int64
```

```
In [84]: sns.countplot(x='Usefulness', order=['useless', '>75%', '25-75%', '<25%'], data=temp_df)
plt.xlabel('Usefulness')
plt.show()
```



### 1.8.2 Positive reviews are found more helpful

As the rating becomes more positive, the reviews become more helpful (and less unhelpful).


```
In [85]: temp_df[temp_df.Score==5].Usefulness.value_counts()
```

```
Out [85]: useless    186743
         >75%      142954
         25-75%    21314
         <25%      12111
         Name: Usefulness, dtype: int64
```

```
In [86]: temp_df[temp_df.Score==2].Usefulness.value_counts()
```

```
Out [86]: useless    10604
         >75%       7423
         25-75%    6693
         <25%      5049
         Name: Usefulness, dtype: int64
```

```
In [87]: sns.countplot(x='Sentiment', hue='Usefulness', order=["positive", "negative"], \
                        hue_order=['>75%', '25-75%', '<25%'], data=temp_df, palette='RdBu')
plt.xlabel('Sentiment')
plt.show()
```



Therefore positive reviews are more helpful.

## 1.9 Word Count

```
In [88]: temp_df["text_word_count"] = temp_df["Text"].apply(lambda text: len(text.split()))
```

```
In [89]: temp_df.head()
```

```
Out [89]:
```

	UserId	HelpfulnessNumerator	HelpfulnessDenominator	\
0	A3SGXH7AUHU8GW	1	1	
1	A1D87F6ZCVE5NK	0	0	
2	ABXLMWJIXXAIN	1	1	
3	A395B0RC6FGVXV	3	3	
4	A1UQRSCLF8GW1T	0	0	

	Summary	Text	\
0	Good Quality Dog Food	I have bought several of the Vitality canned d...	
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...	
2	"Delight" says it all	This is a confection that has been around a fe...	
3	Cough Medicine	If you are looking for the secret ingredient i...	
4	Great taffy	Great taffy at a great price. There was a wid...	

	Score	Sentiment	Usefulness	text_word_count
0	5	positive	>75%	48
1	1	negative	useless	31
2	4	positive	>75%	94
3	2	negative	>75%	41
4	5	positive	useless	27

```
In [90]: temp_df[temp_df.Score==5].text_word_count.median()
```

```
Out [90]: 52.0
```

```
In [91]: temp_df[temp_df.Score==4].text_word_count.median()
```

```
Out [91]: 65.0
```

```
In [92]: temp_df[temp_df.Score==3].text_word_count.median()
```

```
Out [92]: 70.0
```


```
In [93]: temp_df[temp_df.Score==2].text_word_count.median()
```

```
Out [93]: 67.0
```

```
In [94]: temp_df[temp_df.Score==1].text_word_count.median()
```

```
Out [94]: 62.0
```

```
In [104]: sns.boxplot(x='Score',y='text_word_count', data=temp_df, palette='RdBu', showfliers=
plt.show()
```




Observations: 5-star reviews had the lowest median word count (52 words), while 3-star reviews had the largest median word count (70 words).

### 1.9.1 How does word count relate to helpfulness?

The word counts for helpful reviews and not helpful reviews have a similar distribution with the greatest concentration of reviews of approximately 25 words. However, not helpful reviews have a larger concentration of reviews with low word count and helpful reviews have more longer reviews. Helpful reviews have a higher median word count (67 words) than not helpful reviews (54 words).

```
In [105]: sns.violinplot(x='Usefulness', y='text_word_count', order=[">75%", "<25%"], \
                        data=temp_df, palette='RdBu')
plt.ylim(-50, 400)
plt.show()
```



## 1.10 Frequency of reviewers

Using User IDs, one can recognize repeat reviewers. Reviewers that have reviewed over 50 products account for over 5% of all reviews in the database. We will call such reviewers frequent reviewers. (The cutoff choice of 50, as opposed to another choice, seemed to not have a larger impact on the results.) I asked: Does the behavior of frequent reviewers differ from that of infrequent reviewers?

```
In [106]: x = temp_df.UserId.value_counts()
          x.to_dict()
          print("converted Series to dictionary")
```

converted Series to dictionary

```
In [107]: temp_df["reviewer_freq"] = temp_df["UserId"].apply(lambda counts: "Frequent (>50 rev"
                                                             if x[counts]>50 else
```

```
In [108]: temp_df.head()
```

```
Out[108]:
```

	UserId	HelpfulnessNumerator	HelpfulnessDenominator	\
0	A3SGXH7AUHU8GW	1	1	
1	A1D87F6ZCVE5NK	0	0	

2	ABXLMWJIXXAIN	1	1
3	A395BORC6FGVXV	3	3
4	A1UQRSCLF8GW1T	0	0

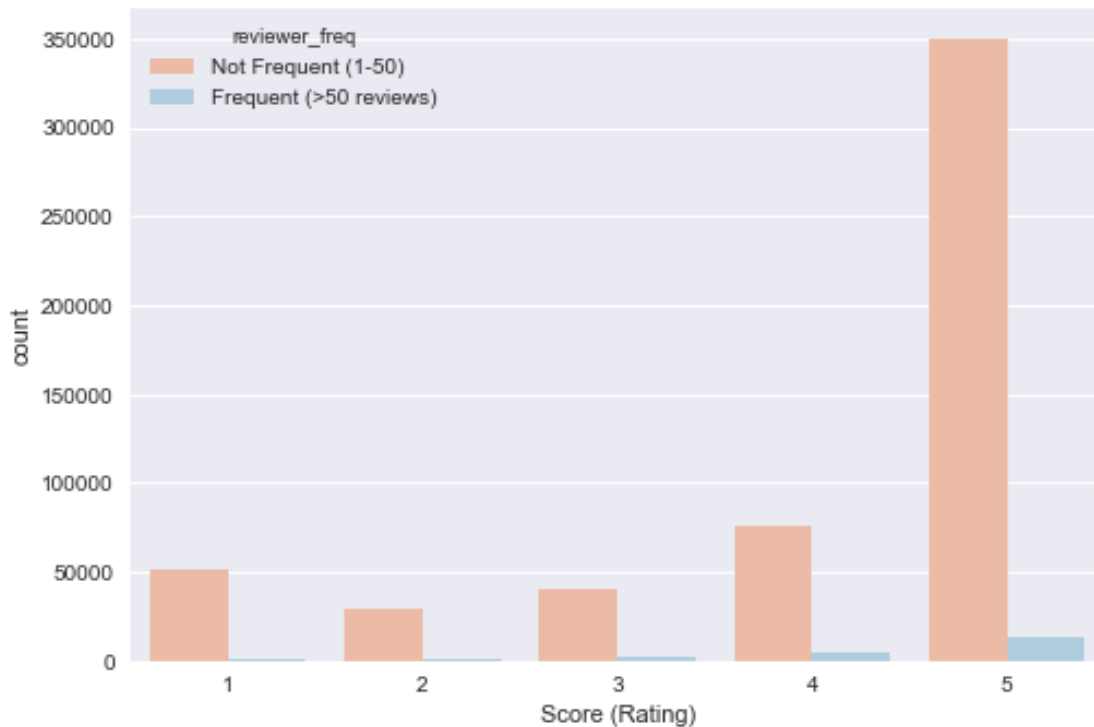
	Summary	Text \
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...
3	Cough Medicine	If you are looking for the secret ingredient i...
4	Great taffy	Great taffy at a great price. There was a wid...

	Score	Sentiment	Usefulness	text_word_count	reviewer_freq
0	5	positive	>75%	48	Not Frequent (1-50)
1	1	negative	useless	31	Not Frequent (1-50)
2	4	positive	>75%	94	Not Frequent (1-50)
3	2	negative	>75%	41	Not Frequent (1-50)
4	5	positive	useless	27	Not Frequent (1-50)

### 1.10.1 Are frequent reviewers more discerning?

The distribution of ratings among frequent reviewers is similar to that of all reviews. However, we can see that frequent reviewers give less 5-star reviews and less 1-star review. Frequent users appear to be more discerning in the sense that they give less extreme reviews than infrequent reviews.

```
In [109]: ax = sns.countplot(x='Score', hue='reviewer_freq', data=temp_df, palette='RdBu')
          ax.set_xlabel('Score (Rating)')
          plt.show()
```



```
In [110]: y = temp_df[temp_df.reviewer_freq=="Frequent (>50 reviews)"].Score.value_counts()
          z = temp_df[temp_df.reviewer_freq=="Not Frequent (1-50)"].Score.value_counts()

          tot_y = y.sum()

          y = (y/tot_y)*100

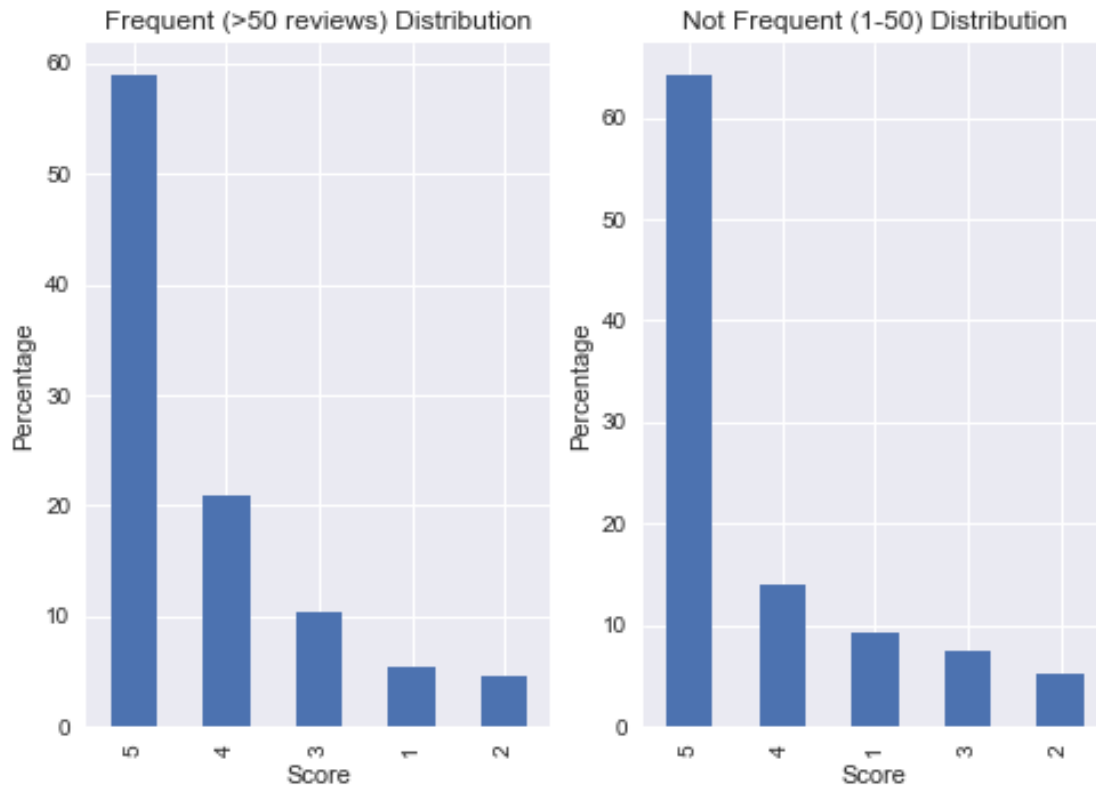
          tot_z = z.sum()

          z = (z/tot_z)*100

          ax1 = plt.subplot(121)
          y.plot(kind="bar",ax=ax1)
          plt.xlabel("Score")
          plt.ylabel("Percentage")
          plt.title("Frequent (>50 reviews) Distribution")

          ax2 = plt.subplot(122)
          z.plot(kind="bar",ax=ax2)
          plt.xlabel("Score")
          plt.ylabel("Percentage")
          plt.title("Not Frequent (1-50) Distribution")
          plt.show()
```

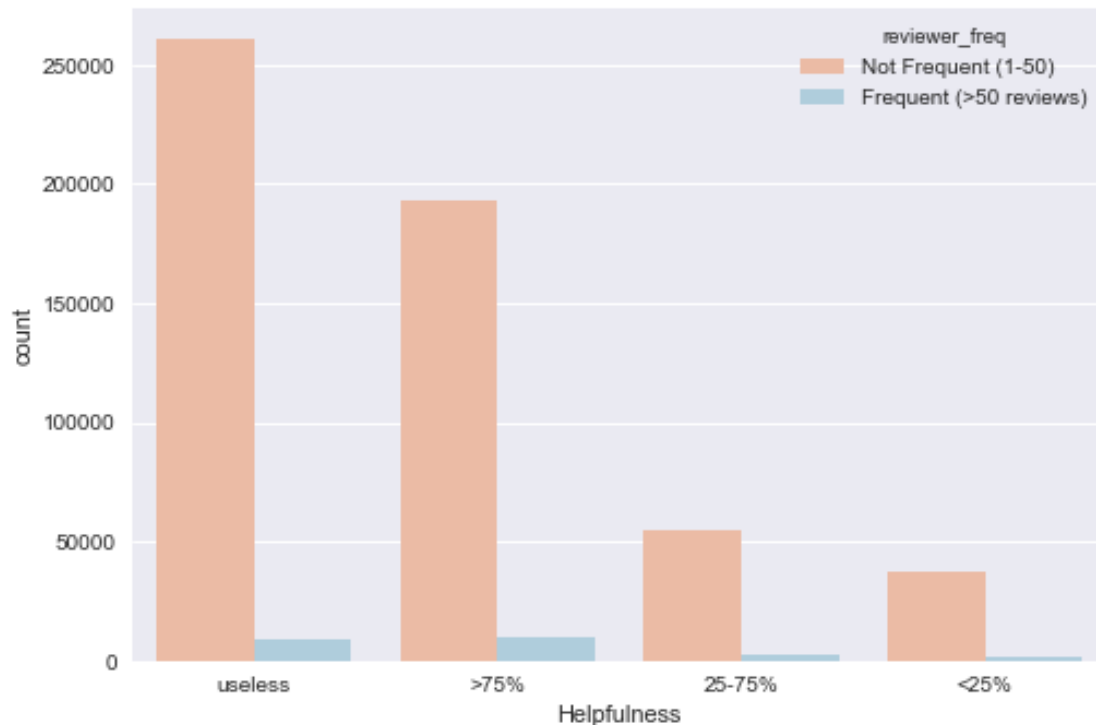




### 1.10.2 Are frequent reviewers more helpful?

The distribution of helpfulness for frequent reviewers is similar to that of all reviews. However, frequent reviewers are more likely to have their review voted on and when voted on, more likely to be voted helpful, and less likely to be unhelpful.

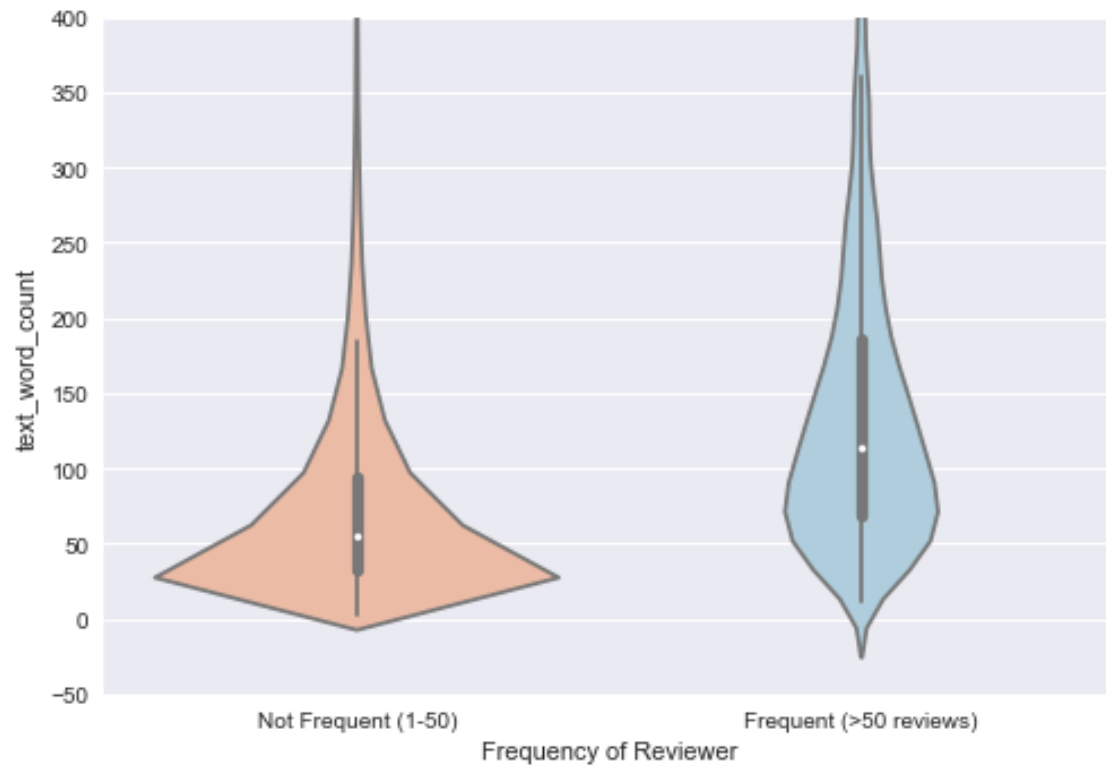
```
In [111]: sns.countplot(x='Usefulness', order=['useless', '>75%', '25-75%', '<25%'], \
                        hue='reviewer_freq', data=temp_df, palette='RdBu')
plt.xlabel('Helpfulness')
plt.show()
```



### 1.10.3 Are frequent reviewers more verbose?

The distributions of word counts for frequent and infrequent reviews shows that infrequent reviewers have a large amount of reviews of low word count. On the other hand, the largest concentration of word count is higher for frequent reviewers than for infrequent reviews. Moreover, the median word count for frequent reviewers is higher than the median for infrequent reviewers.

```
In [112]: sns.violinplot(x='reviewer_freq', y='text_word_count', \
                        data=temp_df, palette='RdBu')
plt.xlabel('Frequency of Reviewer')
plt.ylim(-50, 400)
plt.show()
```



## 1.11 Conclusion